Package ‘bayesCureRateModel’

July 19, 2024

Type Package

Title Bayesian Cure Rate Modeling for Time-to-Event Data

Version 1.1

Date 2024-07-17

Maintainer Panagiotis Papastamoulis <papapast@yahoo.gr>

Description A fully Bayesian approach in order to estimate a general family of cure rate models under the presence of covariates, see Papastamoulis and Milienos (2023) <doi:10.48550/arXiv.2310.06926>. The promotion time can be modelled (a) parametrically using typical distributional assumptions for time to event data (including the Weibull, Exponential, Gompertz, log-Logistic distributions), or (b) semiparametrically using finite mixtures of Gamma distributions. Posterior inference is carried out by constructing a Metropolis-coupled Markov chain Monte Carlo (MCMC) sampler, which combines Gibbs sampling for the latent cure indicators and Metropolis-Hastings steps with Langevin diffusion dynamics for parameter updates. The main MCMC algorithm is embedded within a parallel tempering scheme by considering heated versions of the target posterior distribution.

License GPL-2

URL https://github.com/mqbssppe/Bayesian_cure_rate_model

Imports Rcpp (>= 1.0.12),doParallel, parallel, foreach, mclust, coda, HDInterval, VGAM, calculus, flexsurv

LinkingTo Rcpp, RcppArmadillo

NeedsCompilation yes

Author Panagiotis Papastamoulis [aut, cre] (<https://orcid.org/0000-0001-9468-7613>), Fotios Milienos [aut] (<https://orcid.org/0000-0003-1423-7132>)

Depends R (>= 3.5.0)

R topics documented:

`bayesCureRateModel-package`

Bayesian Cure Rate Modeling for Time-to-Event Data

Description

A fully Bayesian approach in order to estimate a general family of cure rate models under the presence of covariates, see Papastamoulis and Milienos (2023) <doi:10.48550/arXiv.2310.06926>. The promotion time can be modelled (a) parametrically using typical distributional assumptions for time to event data (including the Weibull, Exponential, Gompertz, log-Logistic distributions), or (b) semiparametrically using finite mixtures of Gamma distributions. Posterior inference is carried out by constructing a Metropolis-coupled Markov chain Monte Carlo (MCMC) sampler, which combines Gibbs sampling for the latent cure indicators and Metropolis-Hastings steps with Langevin diffusion dynamics for parameter updates. The main MCMC algorithm is embedded within a parallel tempering scheme by considering heated versions of the target posterior distribution.

The main function of the package is `cure_rate_MC3`. See details for a brief description of the model.

Details

Let $\boldsymbol{y} = (y_1, \ldots, y_n)$ denote the observed data, which correspond to time-to-event data or censoring times. Let also $\boldsymbol{x}_i = (x_{i1}, \ldots, x_{x_{ip}})'$ denote the covariates for subject i, $i = 1, \ldots, n$.

Assuming that the n observations are independent, the observed likelihood is defined as

$$L = L(\boldsymbol{\theta}; \boldsymbol{y}, \boldsymbol{x}) = \prod_{i=1}^{n} f_P(y_i; \boldsymbol{\theta}, \boldsymbol{x}_i)^{\delta_i} S_P(y_i; \boldsymbol{\theta}, \boldsymbol{x}_i)^{1-\delta_i},$$

where $\delta_i = 1$ if the i-th observation corresponds to time-to-event while $\delta_i = 0$ indicates censoring time. The parameter vector $\boldsymbol{\theta}$ is decomposed as

$$\boldsymbol{\theta} = (\boldsymbol{\alpha}', \boldsymbol{\beta}', \gamma, \lambda)$$

where

- $\boldsymbol{\alpha} = (\alpha_1, \ldots, \alpha_d)' \in \mathcal{A}$ are the parameters of the promotion time distribution whose cumulative distribution and density functions are denoted as $F(\cdot, \boldsymbol{\alpha})$ and $f(\cdot, \boldsymbol{\alpha})$, respectively.
- $\boldsymbol{\beta} \in \mathbf{R}^k$ are the regression coefficients with k denoting the number of columns in the design matrix (it may include a constant term or not).
- $\gamma \in \mathbf{R}$
- $\lambda > 0$.

The population survival and density functions are defined as

$$S_P(y; \boldsymbol{\theta}) = \left(1 + \gamma \exp\{\boldsymbol{x}_i \boldsymbol{\beta}'\} c^{\gamma \exp\{\boldsymbol{x}_i \boldsymbol{\beta}'\}} F(y; \boldsymbol{\alpha})^{\lambda}\right)^{-1/\gamma}$$

whereas,

$$f_P(y; \boldsymbol{\theta}) = -\frac{\partial S_P(y; \boldsymbol{\theta})}{\partial y}.$$

Finally, the cure rate is affected through covariates and parameters as follows

$$p_0(\boldsymbol{x}_i; \boldsymbol{\theta}) = \left(1 + \gamma \exp\{\boldsymbol{x}_i \boldsymbol{\beta}'\} c^{\gamma \exp\{\boldsymbol{x}_i \boldsymbol{\beta}'\}}\right)^{-1/\gamma}$$

where $c = e^{e^{-1}}$.

The promotion time distribution can be a member of standard families (currently available are the following: Exponential, Weibull, Gamma, Lomax, Gompertz, log-Logistic) and in this case $\alpha = (\alpha_1, \alpha_2) \in (0, \infty)^2$. Also considered is the Dagum distribution, which has three parameters $(\alpha_1, \alpha_2, \alpha_3) \in (0, \infty)^3$. In case that the previous parametric assumptions are not justified, the promotion time can belong to the more flexible family of finite mixtures of Gamma distributions. For example, assume a mixture of two Gamma distributions of the form

$$f(y; \boldsymbol{\alpha}) = \alpha_5 f_{\mathcal{G}}(y; \alpha_1, \alpha_3) + (1 - \alpha_5) f_{\mathcal{G}}(y; \alpha_2, \alpha_4),$$

where

$$f_{\mathcal{G}}(y; \alpha, \beta) = \frac{\beta^{\alpha}}{\Gamma(\alpha)} y^{\alpha - 1} \exp\{-\beta y\}, y > 0$$

denotes the density of the Gamma distribution with parameters $\alpha > 0$ (shape) and $\beta > 0$ (rate). For the previous model, the parameter vector is

$$\boldsymbol{\alpha} = (\alpha_1, \alpha_2, \alpha_3, \alpha_4, \alpha_5)' \in \mathcal{A}$$

where $\mathcal{A} = (0, \infty)^4 \times (0, 1)$.

More generally, one can fit a mixture of $K > 2$ Gamma distributions. The appropriate model can be selected according to information criteria such as the BIC.

The binary vector $\boldsymbol{I} = (I_1, \ldots, I_n)$ contains the (latent) cure indicators, that is, $I_i = 1$ if the i-th subject is susceptible and $I_i = 0$ if the i-th subject is cured. Δ_0 denotes the subset of $\{1, \ldots, n\}$ containing the censored subjects, whereas $\Delta_1 = \Delta_0^c$ is the (complementary) subset of uncensored subjects. The complete likelihood of the model is

$$L_c(\boldsymbol{\theta}; \boldsymbol{y}, \boldsymbol{I}) = \prod_{i \in \Delta_1} (1 - p_0(\boldsymbol{x}_i, \boldsymbol{\theta})) f_U(y_i; \boldsymbol{\theta}, \boldsymbol{x}_i) \prod_{i \in \Delta_0} p_0(\boldsymbol{x}_i, \boldsymbol{\theta})^{1 - I_i} \{(1 - p_0(\boldsymbol{x}_i, \boldsymbol{\theta})) S_U(y_i; \boldsymbol{\theta}, \boldsymbol{x}_i)\}^{I_i}.$$

f_U and S_U denote the probability density and survival function of the susceptibles, respectively, that is

$$S_U(y_i;\boldsymbol{\theta},\boldsymbol{x}_i)=\frac{S_P(y_i;\boldsymbol{\theta},\boldsymbol{x}_i)-p_0(\boldsymbol{x}_i;\boldsymbol{\theta})}{1-p_0(\boldsymbol{x}_i;\boldsymbol{\theta})}, f_U(y_i;\boldsymbol{\theta},\boldsymbol{x}_i)=\frac{f_P(y_i;\boldsymbol{\theta},\boldsymbol{x}_i)}{1-p_0(\boldsymbol{x}_i;\boldsymbol{\theta})}.$$

Index of help topics:

```
bayesCureRateModel-package
                        Bayesian Cure Rate Modeling for Time-to-Event
                        Data
complete_log_likelihood_general
                        Logarithm of the complete log-likelihood for
                        the general cure rate model.
cure_rate_MC3           Main function of the package
cure_rate_mcmc          The basic MCMC scheme.
log_dagum               PDF and CDF of the Dagum distribution
log_gamma               PDF and CDF of the Gamma distribution
log_gamma_mixture       PDF and CDF of a Gamma mixture distribution
log_gompertz            PDF and CDF of the Gompertz distribution
log_logLogistic         PDF and CDF of the log-Logistic distribution.
log_lomax               PDF and CDF of the Lomax distribution
log_weibull             PDF and CDF of the Weibull distribution
marriage_dataset        Marriage data
plot.bayesCureModel     Plot method
print.bayesCureModel    Print method
summary.bayesCureModel
                        Summary method.
```

Author(s)

Panagiotis Papastamoulis and Fotios S. Milienos

Maintainer: Panagiotis Papastamoulis <papapast@yahoo.gr>

References

Papastamoulis and Milienos (2023). Bayesian inference and cure rate modeling for event history data. arXiv:2310.06926

See Also

cure_rate_MC3

Examples

```
# TOY EXAMPLE (very small numbers... only for CRAN check purposes)
# simulate toy data
set.seed(10)
n = 4
stat = rbinom(n, size = 1, prob = 0.5)
x <- cbind(1, matrix(rnorm(n), n, 1))
```

```
y <- rexp(n)
# run a weibull model with default prior setup
# considering 2 heated chains
fit1 <- cure_rate_MC3(y = y, X = x, Censoring_status = stat,
promotion_time = list(distribution = 'weibull'),
nChains = 2,
nCores = 1,
mcmc_cycles = 3, sweep=2)
# print method
fit1
#  summary method
summary1 <- summary(fit1)

# WARNING: the following parameters
#  mcmc_cycles, nChains
#        should take _larger_ values. E.g. a typical implementation consists of:
#        mcmc_cycles = 15000, nChains = 12


# run a Gamma mixture model with K = 2 components and default prior setup
fit2 <- cure_rate_MC3(y = y, X = x, Censoring_status = stat,
promotion_time = list(
distribution = 'gamma_mixture',
        K = 2),
nChains = 8, nCores = 2,
mcmc_cycles = 10)
summary2 <- summary(fit2)
```

`complete_log_likelihood_general` *Logarithm of the complete log-likelihood for the general cure rate model.*

Description

Compute the logarithm of the complete likelihood, given a realization of the latent binary vector of cure indicators `I_sim` and current values of the model parameters `g`, `lambda`, `b` and promotion time parameters (α) which yield log-density values (one per observation) stored to the vector `log_f` and log-cdf values stored to the vector `log_F`.

Usage

```
complete_log_likelihood_general(y, X, Censoring_status,
g, lambda, log_f, log_F, b, I_sim, alpha)
```

Arguments

- `y` observed data (time-to-event or censored time)
- `X` design matrix. Should contain a column of 1's if the model has a constant term.
- `Censoring_status` binary variables corresponding to time-to-event and censoring.
- `g` The γ parameter of the model (real).
- `lambda` The λ parameter of the model (positive).
- `log_f` A vector containing the natural logarithm of the density function of the promotion time distribution per observation, for the current set of parameters. Its length should be equal to the sample size.
- `log_F` A vector containing the natural logarithm of the cumulative density function of the promotion time distribution per observation, for the current set of parameters. Its length should be equal to the sample size.
- `b` Vector of regression coefficients. Its length should be equal to the number of columns of the design matrix.
- `I_sim` Binary vector of the current value of the latent cured status per observation. Its length should be equal to the sample size. A time-to-event cannot be cured.
- `alpha` A parameter between 0 and 1, corresponding to the temperature of the complete posterior distribution.

Details

The complete likelihood of the model is

$$L_c(\boldsymbol{\theta}; \boldsymbol{y}, \boldsymbol{I}) = \prod_{i\in\Delta_1}(1-p_0(\boldsymbol{x}_i,\boldsymbol{\theta}))f_U(y_i;\boldsymbol{\theta},\boldsymbol{x}_i)\prod_{i\in\Delta_0}p_0(\boldsymbol{x}_i,\boldsymbol{\theta})^{1-I_i}\{(1-p_0(\boldsymbol{x}_i,\boldsymbol{\theta}))S_U(y_i;\boldsymbol{\theta},\boldsymbol{x}_i)\}^{I_i}.$$

f_U and S_U denote the probability density and survival function of the susceptibles, respectively, that is

$$S_U(y_i;\boldsymbol{\theta},\boldsymbol{x}_i) = \frac{S_P(y_i;\boldsymbol{\theta},\boldsymbol{x}_i) - p_0(\boldsymbol{x}_i;\boldsymbol{\theta})}{1-p_0(\boldsymbol{x}_i;\boldsymbol{\theta})}, f_U(y_i;\boldsymbol{\theta},\boldsymbol{x}_i) = \frac{f_P(y_i;\boldsymbol{\theta},\boldsymbol{x}_i)}{1-p_0(\boldsymbol{x}_i;\boldsymbol{\theta})}.$$

Value

A list with the following entries

- `cll` the complete log-likelihood for the current parameter values.
- `logS` Vector of logS values (one for each observation).
- `logP0` Vector of logP0 values (one for each observation).

Author(s)

Panagiotis Papastamoulis

References

Papastamoulis and Milienos (2023). Bayesian inference and cure rate modeling for event history data. arXiv:2310.06926.

Examples

```
# simulate toy data
set.seed(1)
n = 4
stat = rbinom(n, size = 1, prob = 0.5)
x <- cbind(1, matrix(rnorm(n), n, 1))
y <- rexp(n)
lw <- log_weibull(y, a1 = 1, a2 = 1, c_under = 1e-9)
# compute complete log-likelihood
complete_log_likelihood_general(y = y, X = x,
Censoring_status = stat,
g = 1, lambda = 1,
log_f = lw$log_f, log_F = lw$log_F,
b = c(-0.5,0.5),
I_sim = stat, alpha = 1)
```

`cure_rate_MC3` *Main function of the package*

Description

Runs a Metropolis Coupled MCMC (MC^3) sampler in order to estimate the joint posterior distribution of the model.

Usage

```
cure_rate_MC3(y, X, Censoring_status, nChains = 12, mcmc_cycles = 15000,
alpha = NULL,nCores = 8, sweep = 5, mu_g = 1, s2_g = 1,
a_l = 2.1, b_l = 1.1, mu_b = rep(0, dim(X)[2]),
Sigma = 100 * diag(dim(X)[2]), g_prop_sd = 0.045,
lambda_prop_scale = 0.03, b_prop_sd = rep(0.022, dim(X)[2]),
initialValues = NULL, plot = TRUE, adjust_scales = FALSE,
verbose = FALSE, tau_mala = 1.5e-05, mala = 0.15,
promotion_time = list(distribution = "weibull",
prior_parameters = matrix(rep(c(2.1, 1.1), 2), byrow = TRUE, 2, 2),
prop_scale = c(0.1, 0.2)), single_MH_in_f = 0.2)
```

Arguments

`y` Observed data, that is, a vector of length n with positive entries.

`X` Design matrix with $k > 1$ columns. Should contain a column of 1's if the model has a constant term. Should be a matrix with dimension $n \times k$.

`Censoring_status` A vector $\boldsymbol{\delta} = (\delta_1, \ldots, \delta_n)$ of binary variables corresponding to censoring indicators. The i-th observation is treated as a time-to-event if $\delta_i = 1$ or as a censoring time otherwise ($\delta_i = 0$).

nChains — Positive integer corresponding to the number of heated chains in the MC^3 scheme.

mcmc_cycles — Length of the generated MCMC sample. Default value: 15000. Note that each MCMC cycle consists of sweep (see below) usual MCMC iterations.

alpha — A decreasing sequence in $[1, 0)$ of nChains temperatures (or heat values). The first value should always be equal to 1, which corresponds to the target posterior distribution (that is, the first chain).

nCores — The number of cores used for parallel processing.

sweep — The number of usual MCMC iterations per MCMC cycle. Default value: 10.

mu_g — Parameter a_γ of the prior distribution of γ.

s2_g — Parameter b_γ of the prior distribution of γ.

a_l — Shape parameter a_λ of the Inverse Gamma prior distribution of λ.

b_l — Scale parameter b_λ of the Inverse Gamma prior distribution of λ.

mu_b — Mean ($\boldsymbol{\mu}$) of the multivariate normal prior distribution of regression coefficients. Should be a vector whose length is equal to k, i.e. the number of columns of the design matrix X. Default value: rep(0, k).

Sigma — Covariance matrix of the multivariate normal prior distribution of regression coefficients.

g_prop_sd — The scale of the proposal distribution for single-site updates of the γ parameter.

lambda_prop_scale — The scale of the proposal distribution for single-site updates of the λ parameter.

b_prop_sd — The scale of the proposal distribution for the update of the β parameter (regression coefficients).

initialValues — A list of initial values for each parameter (optional).

plot — Plot MCMC sample on the run. Default: TRUE.

adjust_scales — Boolean. If TRUE the MCMC sampler runs an initial phase of a small number of iterations in order to tune the scale of the proposal distributions in the Metropolis-Hastings steps.

verbose — Print progress on the terminal if TRUE.

tau_mala — Scale of the Metropolis adjusted Langevin diffussion proposal distribution.

mala — A number between $[0, 1]$ indicating the proportion of times the sampler attempts a MALA proposal. Thus, the probability of attempting a typical Metropolis-Hastings move is equal to 1 - mala.

promotion_time — A list with details indicating the parametric family of distribution describing the promotion time and corresponding prior distributions. See 'details'.

single_MH_in_f — The probability for attempting a series of single site updates in the typical Metropolis-Hastings move. Otherwise, with probability 1 - single_MH_in_f a Metropolis-Hastings move will attempt to update all continuous parameters simultaneously. It only makes sense when mala < 1.

Details

It is advised to scale all continuous explanatory variables in the design matrix, so their sample mean and standard deviations are equal to 0 and 1, respectively. The `promotion_time` argument should be a list containing the following entries

`distribution` Character string specifying the family of distributions $\{F(\cdot;\boldsymbol{\alpha});\boldsymbol{\alpha} \in \mathcal{A}\}$ describing the promotion time.

`prior_parameters` Values of hyper-parameters in the prior distribution of the parameters $\boldsymbol{\alpha}$.

`prop_scale` The scale of the proposal distributions for each parameter in $\boldsymbol{\alpha}$.

`dirichlet_concentration_parameter` Relevant only in the case of the `'gamma_mixture'`. Positive scalar (typically, set to 1) determining the (common) concentration parameter of the Dirichlet prior distribution of mixing proportions.

The `distribution` entry should be one of the following: `'exponential'`, `'weibull'`, `'gamma'`, `'logLogistic'`, `'gompertz'`, `'lomax'`, `'dagum'`, `'gamma_mixture'`.

The joint prior distribution of $\boldsymbol{\alpha} = (\alpha_1, \ldots, \alpha_d)$ factorizes into products of inverse Gamma distributions for all (positive) parameters of F. Moreover, in the case of `'gamma_mixture'`, the joint prior also consists of another term to the Dirichlet prior distribution on the mixing proportions.

The `prop_scale` argument should be a vector with length equal to the length of vector d (number of elements in $\boldsymbol{\alpha}$), containing (positive) numbers which correspond to the scale of the proposal distribution. Note that these scale parameters are used only as initial values in case where `adjust_scales = TRUE`.

Value

An object of class `bayesCureModel`, i.e. a list with the following entries

`mcmc_sample` Object of class `mcmc` (see the **coda** package), containing the generated MCMC sample for the target chain. The column names correspond to

- `g_mcmc` Sampled values for parameter γ
- `lambda_mcmc` Sampled values for parameter λ
- `alpha1_mcmc`**...** Sampled values for parameter α_1... of the promotion time distribution $F(\cdot;\alpha_1,\ldots,\alpha_d)$. The subsequent $d-1$ columns contain the sampled values for all remaining parameters, $\alpha_2, ..., \ldots, \alpha_d$, where d depens on the family used in `promotion_time`.
- `b0_mcmc` Sampled values for the constant term of the regression (present only in the case where the design matrix `X` contains a column of 1s).
- `b1_mcmc`**...** Sampled values for the regression coefficient for the first explanatory variable, and similar for all subsequent columns.

`complete_log_likelihood`
The complete log-likelihood for the target chain.

`all_cll_values`
The complete log-likelihood for all chains

`latent_status_censored`
A data frame with the simulated binary latent status for each censored item.

swap_accept_rate
: the acceptance rate of proposed swappings between adjacent MCMC chains.

input_data_and_model_prior
: the input data and specification of the prior parameters.

log_posterior
: the logarithm of the posterior distribution, up to a normalizing constant.

map_estimate
: The Maximum A Posterior estimate of parameters

BIC
: Bayesian Information Criterion.

AIC
: Akaike Information Criterion.

Note

The core function is cure_rate_mcmc.

Author(s)

Panagiotis Papastamoulis

References

Papastamoulis and Milienos (2023). Bayesian inference and cure rate modeling for event history data. arXiv:2310.06926

See Also

cure_rate_mcmc

Examples

```
# simulate toy data just for cran-check purposes
set.seed(10)
        n = 4
        stat = rbinom(n, size = 1, prob = 0.5)
        x <- cbind(1, matrix(rnorm(2*n), n, 2))
        y <- rexp(n)
fit1 <- cure_rate_MC3(y = y, X = x, Censoring_status = stat,
promotion_time = list(distribution = 'weibull'),
nChains = 2, nCores = 1,
mcmc_cycles = 3, sweep = 2)
```

`cure_rate_mcmc` *The basic MCMC scheme.*

Description

This is core MCMC function. The continuous parameters of the model are updated using (a) single-site Metropolis-Hastings steps and (b) a Metropolis adjusted Langevin diffusion step. The binary latent variables of the model (cured status per censored observation) are updated according to a Gibbs step. This function is embedded to the main function of the package `cure_rate_MC3` which runs parallel tempered MCMC chains.

Usage

```
cure_rate_mcmc(y, X, Censoring_status, m, alpha = 1, mu_g = 1, s2_g = 1,
a_l = 2.1, b_l = 1.1, promotion_time = list(distribution = "weibull",
prior_parameters = matrix(rep(c(2.1, 1.1), 2), byrow = TRUE, 2, 2),
prop_scale = c(0.2, 0.03)), mu_b = NULL, Sigma = NULL, g_prop_sd = 0.045,
lambda_prop_scale = 0.03, b_prop_sd = NULL, initialValues = NULL,
plot = FALSE, verbose = FALSE, tau_mala = 1.5e-05, mala = 0.15,
single_MH_in_f = 0.5)
```

Arguments

`y`	observed data (time-to-event or censored time)
`X`	design matrix. Should contain a column of 1's if the model has a constant term.
`Censoring_status`	binary variables corresponding to time-to-event and censoring.
`m`	number of MCMC iterations.
`alpha`	A value between 0 and 1, corresponding to the temperature of the complete posterior distribution. The target posterior distribution corresponds to `alpha = 1`.
`mu_g`	Parameter a_γ of the prior distribution of γ.
`s2_g`	Parameter b_γ of the prior distribution of γ.
`a_l`	Shape parameter a_λ of the Inverse Gamma prior distribution of λ.
`b_l`	Scale parameter b_λ of the Inverse Gamma prior distribution of λ.
`promotion_time`	A list containing the specification of the promotion time distribution. See 'details'.
`mu_b`	Mean μ of the multivariate normal prior distribution of regression coefficients. Should be a vector whose length is equal to the number of columns of the design matrix `X`.
`Sigma`	Covariance matrix of the multivariate normal prior distribution of regression coefficients.
`g_prop_sd`	The scale of the proposal distribution for single-site updates of the γ parameter.

lambda_prop_scale
The scale of the proposal distribution for single-site updates of the λ parameter.

b_prop_sd
The scale of the proposal distribution for the update of the β parameter (regression coefficients).

initialValues
A list of initial values for each parameter (optional).

plot
Boolean for plotting on the run.

verbose
Boolean for printing progress on the run.

tau_mala
scale of the MALA proposal.

mala
Propability of attempting a MALA step. Otherwise, a simple MH move is attempted.

single_MH_in_f
Probability of attempting a single-site MH move in the basic Metropolis-Hastings step. Otherwise, a joint update is attempted.

Value

A list containing the following entries

mcmc_sample
The sampled MCMC values per parameter. See 'note'.

complete_log_likelihood
Logarithm of the complete likelihood per MCMC iteration.

acceptance_rates
The acceptance rate per move.

latent_status_censored
The MCMC sample of the latent status per censored observation.

log_prior_density
Logarithm of the prior density per MCMC iteration.

Note

In the case where the promotion time distribution is a Gamma mixture model, the mixing proportions $w_1, \ldots, w_K$ are reparameterized according to the following transformation

$$w_j = \frac{\rho_j}{\sum_{i=1}^{K} \rho_i}, j = 1, \ldots, K$$

where $\rho_i > 0$ for $i = 1, \ldots, K-1$ and $\rho_K = 1$. The sampler returns the parameters $\rho_1, \ldots, \rho_{K-1}$, not the mixing proportions.

Author(s)

Panagiotis Papastamoulis

References

Papastamoulis and Milienos (2023). Bayesian inference and cure rate modeling for event history data. arXiv:2310.06926

See Also

cure_rate_MC3

Examples

```
# simulate toy data just for cran-check purposes
        set.seed(1)
        n = 10
        stat = rbinom(n, size = 1, prob = 0.5)
        x <- cbind(1, matrix(rnorm(2*n), n, 2))
        y <- rexp(n)
# run a weibull model (including const. term)
# for m = 10 mcmc iterations
        fit1 <- cure_rate_mcmc(y = y, X = x, Censoring_status = stat,
               plot = FALSE,
                promotion_time = list(distribution = 'weibull',
                        prior_parameters = matrix(rep(c(2.1, 1.1), 2),
                                                byrow = TRUE, 2, 2),
                        prop_scale = c(0.1, 0.1)
                ),
                m = 10)
# the generated mcmc sampled values
fit1$mcmc_sample
```

---

log_dagum — *PDF and CDF of the Dagum distribution*

Description

The Dagum distribution as evaluated at the **VGAM** package.

Usage

```
log_dagum(y, a1, a2, a3, c_under = 1e-09)
```

Arguments

y	observed data
a1	scale parameter
a2	shape1.a parameter
a3	shape2.p parameter
c_under	A small positive value corresponding to the underflow threshold, e.g. c_under = 1e-9.

Details

The Dagum distribution is a special case of the 4-parameter generalized beta II distribution.

Value

A list containing the following entries

`log_f`	natural logarithm of the pdf, evaluated at each datapoint.
`log_F`	natural logarithm of the CDF, evaluated at each datapoint.

Author(s)

Panagiotis Papastamoulis

References

Thomas W. Yee (2015). Vector Generalized Linear and Additive Models: With an Implementation in R. New York, USA: Springer.

See Also

`ddagum`

Examples

```
log_dagum(y = 1:10, a1 = 1, a2 = 1, a3 = 1, c_under = 1e-9)
```

`log_gamma` *PDF and CDF of the Gamma distribution*

Description

Computes the pdf and cdf of the Gamma distribution.

Usage

```
log_gamma(y, a1, a2, c_under = 1e-09)
```

Arguments

`y`	observed data
`a1`	shape parameter
`a2`	rate parameter
`c_under`	A small positive value corresponding to the underflow threshold, e.g. c_under = 1e-9.

Value

A list containing the following entries

`log_f`	natural logarithm of the pdf, evaluated at each datapoint.
`log_F`	natural logarithm of the CDF, evaluated at each datapoint.

Author(s)

Panagiotis Papastamoulis

See Also

dgamma

Examples

```
log_gamma(y = 1:10, a1 = 1, a2 = 1, c_under = 1e-9)
```

`log_gamma_mixture` *PDF and CDF of a Gamma mixture distribution*

Description

Computes the logarithm of the probability density function and cumulative density function per observation for each observation under a Gamma mixture model.

Usage

```
log_gamma_mixture(y, a1, a2, p, c_under = 1e-09)
```

Arguments

`y`	observed data
`a1`	vector containing the shape parameters of each Gamma mixture component
`a2`	vector containing the rate parameters of each Gamma mixture component
`p`	vector of mixing proportions
`c_under`	threshold for underflows.

Value

A list containing the following entries

`log_f`	natural logarithm of the pdf, evaluated at each datapoint.
`log_F`	natural logarithm of the CDF, evaluated at each datapoint.

Author(s)

Panagiotis Papastamoulis

Examples

```
y <- runif(10)
a1 <- c(1,2)
a2 <- c(1,1)
p <- c(0.9,0.1)
log_gamma_mixture(y, a1, a2, p)
```

log_gompertz *PDF and CDF of the Gompertz distribution*

Description

The Gompertz distribution as evaluated at the **flexsurv** package.

Usage

```
log_gompertz(y, a1, a2, c_under = 1e-09)
```

Arguments

y	observed data
a1	shape parameter
a2	rate parameter
c_under	A small positive value corresponding to the underflow threshold, e.g. c_under = 1e-9.

Value

A list containing the following entries

log_f	natural logarithm of the pdf, evaluated at each datapoint.
log_F	natural logarithm of the CDF, evaluated at each datapoint.

Author(s)

Panagiotis Papastamoulis

References

Christopher Jackson (2016). flexsurv: A Platform for Parametric Survival Modeling in R. Journal of Statistical Software, 70(8), 1-33. doi:10.18637/jss.v070.i08

See Also

dgompertz

Examples

```
log_gompertz(y = 1:10, a1 = 1, a2 = 1, c_under = 1e-9)
```

log_logLogistic — *PDF and CDF of the log-Logistic distribution.*

Description

The log-Logistic distribution as evaluated at the **flexsurv** package.

Usage

```
log_logLogistic(y, a1, a2, c_under = 1e-09)
```

Arguments

`y`	observed data
`a1`	shape parameter
`a2`	scale parameter
`c_under`	A small positive value corresponding to the underflow threshold, e.g. c_under = 1e-9.

Details

The log-logistic distribution is the probability distribution of a random variable whose logarithm has a logistic distribution.

Value

A list containing the following entries

`log_f`	natural logarithm of the pdf, evaluated at each datapoint.
`log_F`	natural logarithm of the CDF, evaluated at each datapoint.

Author(s)

Panagiotis Papastamoulis

References

Christopher Jackson (2016). flexsurv: A Platform for Parametric Survival Modeling in R. Journal of Statistical Software, 70(8), 1-33. doi:10.18637/jss.v070.i08

See Also

`dllogis`

Examples

```
log_logLogistic(y = 1:10, a1 = 1, a2 = 1, c_under = 1e-9)
```

`log_lomax` *PDF and CDF of the Lomax distribution*

Description

The Lomax distribution as evaluated at the **VGAM** package.

Usage

```
log_lomax(y, a1, a2, c_under = 1e-09)
```

Arguments

`y`	observed data
`a1`	scale parameter
`a2`	shape parameter
`c_under`	A small positive value corresponding to the underflow threshold, e.g. c_under = 1e-9.

Details

The Lomax distribution is a special case of the 4-parameter generalized beta II distribution.

Value

A list containing the following entries

`log_f`	natural logarithm of the pdf, evaluated at each datapoint.
`log_F`	natural logarithm of the CDF, evaluated at each datapoint.

Author(s)

Panagiotis Papastamoulis

References

Thomas W. Yee (2015). Vector Generalized Linear and Additive Models: With an Implementation in R. New York, USA: Springer.

See Also

`dlomax`

Examples

```
log_lomax(y = 1:10, a1 = 1, a2 = 1, c_under = 1e-9)
```

log_weibull *PDF and CDF of the Weibull distribution*

Description

Computes the log pdf and cdf of the weibull distribution.

Usage

```
log_weibull(y, a1, a2, c_under)
```

Arguments

y	observed data
a1	shape parameter
a2	rate parameter
c_under	A small positive value corresponding to the underflow threshold, e.g. c_under = 1e-9.

Value

A list containing the following entries

log_f	natural logarithm of the pdf, evaluated at each datapoint.
log_F	natural logarithm of the CDF, evaluated at each datapoint.

Author(s)

Panagiotis Papastamoulis

Examples

```
log_weibull(y = 1:10, a1 = 1, a2 = 1, c_under = 1e-9)
```

marriage_dataset *Marriage data*

Description

The variable of interest (time) corresponds to the duration (in years) of first marriage for 1500 individuals. The available covariates are:

- `age` age of respondent (in years) at the time of fist marriage.
- `kids` factor: whether there were kids during the first marriage (1) or not (0).
- `race` race of respondent decoded as: black (1), hispanic (2) and non-black/non-hispanic (4).

Among the 1500 observations, there are 1018 censoring times (censoring = 0) and 482 divorces (censoring = 1). Source: National Longitudinal Survey of Youth 1997 (NLSY97).

Usage

```
data(marriage_dataset)
```

Format

Time-to-event data.

References

Bureau of Labor Statistics, U.S. Department of Labor. National Longitudinal Survey of Youth 1997 cohort, 1997-2022 (rounds 1-20). Produced and distributed by the Center for Human Resource Research (CHRR), The Ohio State University. Columbus, OH: 2023.

plot.bayesCureModel *Plot method*

Description

Plots and computes HDIs.

Usage

```
## S3 method for class 'bayesCureModel'
plot(x, burn = NULL, alpha = 0.05, gamma_mix = TRUE,
K_gamma = 5, plot_graphs = TRUE, bw = "nrd0", what = NULL, p_cured_output = NULL,
index_of_main_mode = NULL,...)
```

Arguments

- `x` — An object of class `bayesCureModel`
- `burn` — Number of iterations to discard as burn-in period.
- `alpha` — A value between 0 and 1 in order to compute the 1-α Highest Posterior Density regions.
- `gamma_mix` — Boolean. If TRUE, the density of the marginal posterior distribution of the γ parameter is estimated from the sampled MCMC values by fitting a normal mixture model.
- `K_gamma` — Used only when `gamma_mix = TRUE` and corresponds to the number of normal mixture components used to estimate the marginal posterior density of the γ parameter.
- `plot_graphs` — Boolean, if FALSE only HDIs are computed.
- `bw` — bandwidth
- `what` — Integer or a character string with possible values equal to `'cured_prob'` or `'survival'`. An integer entry indicates which parameter should be plotted. If set to `NULL` (default), all parameters are plotted one by one. If set to `'cured_prob'` or `'survival'` the estimated cured probability or survival function is plotted, conditional on a set of covariates defined in the `p_cured_output` argument.
- `p_cured_output` — Optional argument (list) which is required only when `what = 'cured_prob'` or `what = 'survival'`. It is returned by the `summary.bayesCureRateModel` function.
- `index_of_main_mode` — If NULL (default), all modes are plotted. Otherwise, it is a subset of the retained MCMC iterations in order to identify the main mode of the posterior distribution, as returned by the `index_of_main_mode` value of the `summary.bayesCureRateModel` function.
- `...` — arguments passed by other methods.

Value

The function plots graphic output on the plot device if `plot_graphs = TRUE`. Furthermore, a list of $100(1-\alpha)\%$ Highest Density Intervals per parameter is returned.

Author(s)

Panagiotis Papastamoulis

Examples

```
# simulate toy data just for cran-check purposes
        set.seed(1)
        n = 4
        stat = rbinom(n, size = 1, prob = 0.5)
        # simulate design matrix
        # first column consists of 1s (const)
        # and second and third column contain
```

```
        # the values of two covariates
        x <- cbind(1, matrix(rnorm(2*n), n, 2))
        colnames(x) <- c('const', 'x1', 'x2')
        y <- rexp(n)
fit1 <- cure_rate_MC3(y = y, X = x, Censoring_status = stat,
promotion_time = list(distribution = 'exponential'),
nChains = 2, nCores = 1,
mcmc_cycles = 3, sweep = 2)
# plot the marginal posterior distribution of the first parameter in returned mcmc output
plot(fit1, what = 1, burn = 0)
# using 'cured_prob'

#compute cured probability for two individuals with
# x1 = 0.2 and x2 = -1
# and
# x1 = -1 and x2 = 0
covariate_levels1 <- rbind(c(1,0.2,-1), c(1,-1,0))
summary1 <- summary(fit1, covariate_levels = covariate_levels1, burn = 0)
plot(fit1, what='cured_prob', p_cured_output = summary1$p_cured_output,
  ylim = c(0,1))
```

`print.bayesCureModel` *Print method*

Description

This function prints a summary of objects returned by the `cure_rate_MC3` function.

Usage

```
## S3 method for class 'bayesCureModel'
print(x, ...)
```

Arguments

`x`	An object of class `bayesCureModel`, which is returned by the `cure_rate_MC3` function.
`...`	ignored.

Details

The function prints some basic information for a `cure_rate_MC3`, such as the MAP estimate of model parameters and the value of Bayesian information criterion.

Value

No return value, called for side effects.

Author(s)

Panagiotis Papastamoulis

`summary.bayesCureModel` *Summary method.*

Description

This function produces all summaries after fitting a cure rate model.

Usage

```
## S3 method for class 'bayesCureModel'
summary(object, burn = NULL, gamma_mix = TRUE,
K_gamma = 3,  K_max = 3, fdr = 0.1,
covariate_levels = NULL, yRange = NULL, alpha = 0.1, ...)
```

Arguments

`object`	An object of class `bayesCureModel`
`burn`	Positive integer corresponding to the number of mcmc iterations to discard as burn-in period
`gamma_mix`	Boolean. If TRUE, the density of the marginal posterior distribution of the γ parameter is estimated from the sampled MCMC values by fitting a normal mixture model.
`K_gamma`	Used only when `gamma_mix = TRUE` and corresponds to the number of normal mixture components used to estimate the marginal posterior density of the γ parameter.
`K_max`	Maximum number of components in order to cluster the (univariate) values of the joint posterior distribution across the MCMC run. Used to identify the main mode of the posterior distribution. See details.
`fdr`	The target value for controlling the False Discovery Rate when classifying subjects as cured or not.
`covariate_levels`	Optional levels for the covariates. It is only required when the user wishes to obtain a vector with the estimated posterior cured probabilities for a given combination of covariates. Include the value "1" in the case where the model contains constant term.
`yRange`	Optional range (a vector of two non-negative values) for computing the sequence of posterior probabilities for the given values in `covariate_levels`.

`alpha` Scalar between 0 and 1 corresponding to 1 - confidencel level for computing Highest Density Intervals. If set to NULL, the confidence intervals are not computed.

`...` ignored.

Details

The values of the posterior draws are clustered according to a (univariate) normal mixture model, and the main mode corresponds to the cluster with the largest mean. The maximum number of mixture components corresponds to the `K_max` argument. The **mclust** library is used for this purpose. The inference for the latent cure status of each (censored) observation is based on the MCMC draws corresponding to the main mode of the posterior distribution. The FDR is controlled according to the technique proposed in Papastamoulis and Rattray (2018).

In case where `covariate_levels` is set to `TRUE`, the `summary` function also returns a list named `p_cured_output` with the following entries

mcmc It is returned only in the case where the argument `covariate_values` is not `NULL`. A vector of posterior cured probabilities for the given values in `covariate_values`, per retained MCMC draw.

map It is returned only in the case where the argument `covariate_values` is not `NULL`. The cured probabilities computed at the MAP estimate of the parameters, for the given values `covariate_values`.

tau_values tau values

covariate_levels covariate levels

index_of_main_mode the subset of MCMC draws allocated to the main mode of the posterior distribution.

Value

A list with the following entries

`map_estimate` Maximum A Posteriori (MAP) estimate of the parameters of the model.

`highest_density_intervals` Highest Density Interval per parameter

`latent_cured_status` Estimated posterior probabilities of the latent cure status per censored subject.

`cured_at_given_FDR` Classification as cured or not, at given FDR level.

`p_cured_output` It is returned only in the case where the argument `covariate_values` is not `NULL`. See details.

`main_mode_index` The retained MCMC iterations which correspond to the main mode of the posterior distribution.

Author(s)

Panagiotis Papastamoulis

References

Papastamoulis and Milienos (2023). Bayesian inference and cure rate modeling for event history data. arXiv:2310.06926.

Papastamoulis and Rattray (2018). A Bayesian Model Selection Approach for Identifying Differentially Expressed Transcripts from RNA Sequencing Data, Journal of the Royal Statistical Society Series C: Applied Statistics, Volume 67, Issue 1.

Scrucca L, Fraley C, Murphy TB, Raftery AE (2023). Model-Based Clustering, Classification, and Density Estimation Using mclust in R. Chapman and Hall/CRC. ISBN 978-1032234953

See Also

cure_rate_MC3

Examples

```
# simulate toy data just for cran-check purposes
        set.seed(1)
        n = 4
        stat = rbinom(n, size = 1, prob = 0.5)
        x <- cbind(1, matrix(rnorm(2*n), n, 2))
        y <- rexp(n)
fit1 <- cure_rate_MC3(y = y, X = x, Censoring_status = stat,
promotion_time = list(distribution = 'exponential'),
nChains = 2, nCores = 1,
mcmc_cycles = 3, sweep = 2)
mySummary <- summary(fit1, burn = 0)
```

Index